cloudera

Petabytes of Data - How to Use R at Scale

Jeff Fletcher

AGENDA

- WHAT PROBLEM ARE WE SOLVING?
- HOW IS IT SOLVED OUTSIDE OF R?
- COMPONENTS OVERVIEW
- HOW DOES THIS WORK WITH R?
- USEFUL THINGS TO KNOW
- GETTING STARTED

WHAT PROBLEM ARE WE SOLVING?

WHAT PROBLEM ARE WE SOLVING?

Your MacBook, while very cool and shiny is too small for the enormous amounts of data used by actual enterprise companies.

WHAT PROBLEM ARE WE SOLVING?

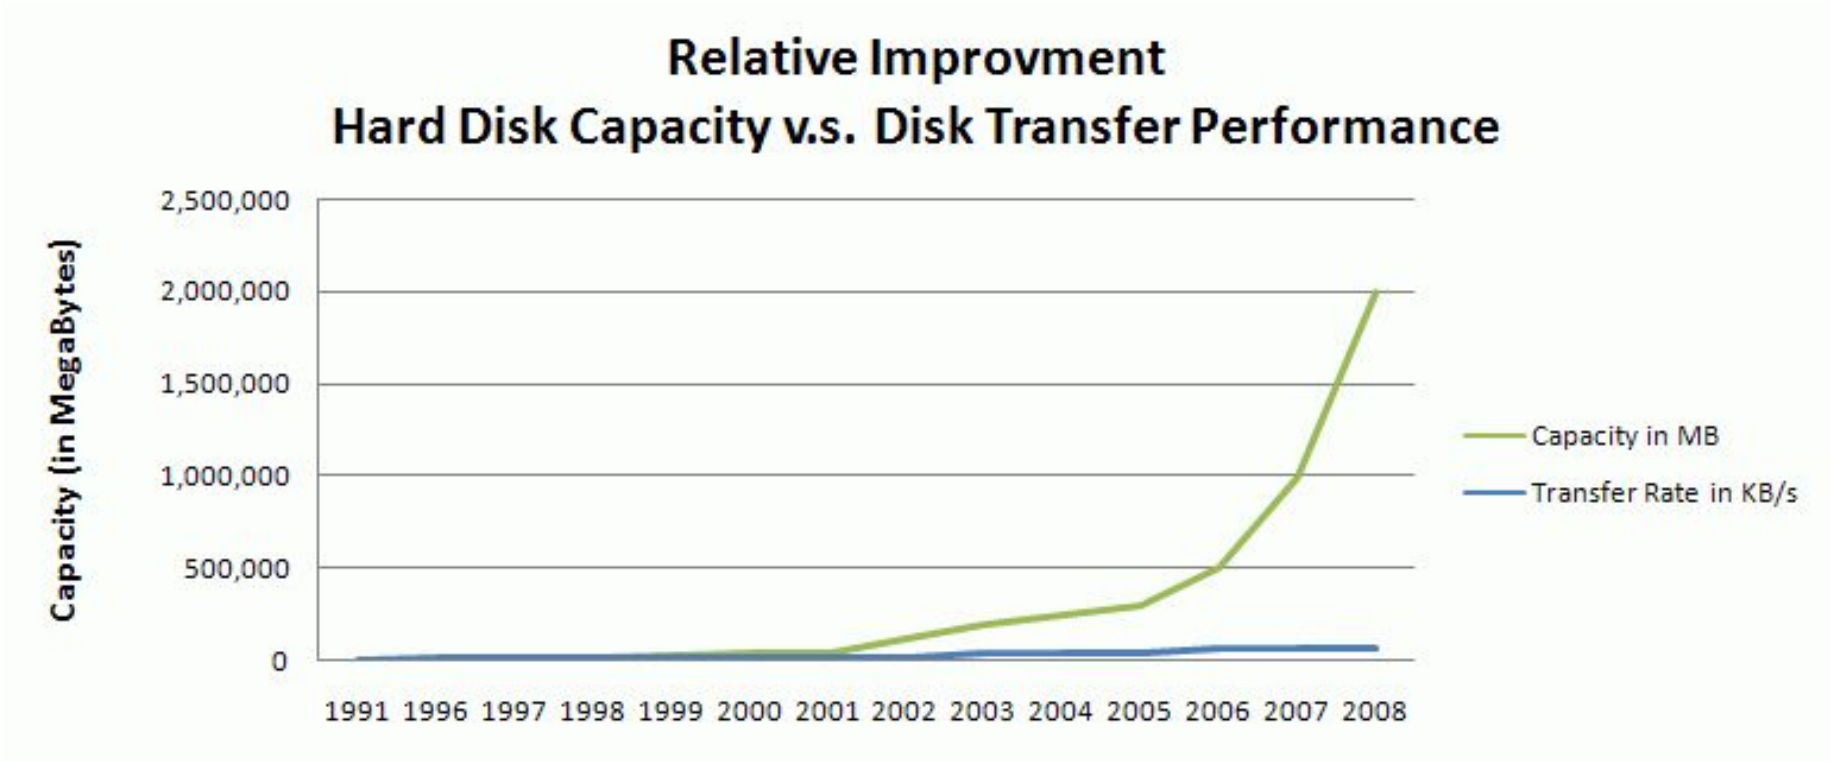

WHAT PROBLEM ARE WE SOLVING?

"It's not who has the best algorithms that wins. It's who has the most data." [Banko and Brill, 2001]

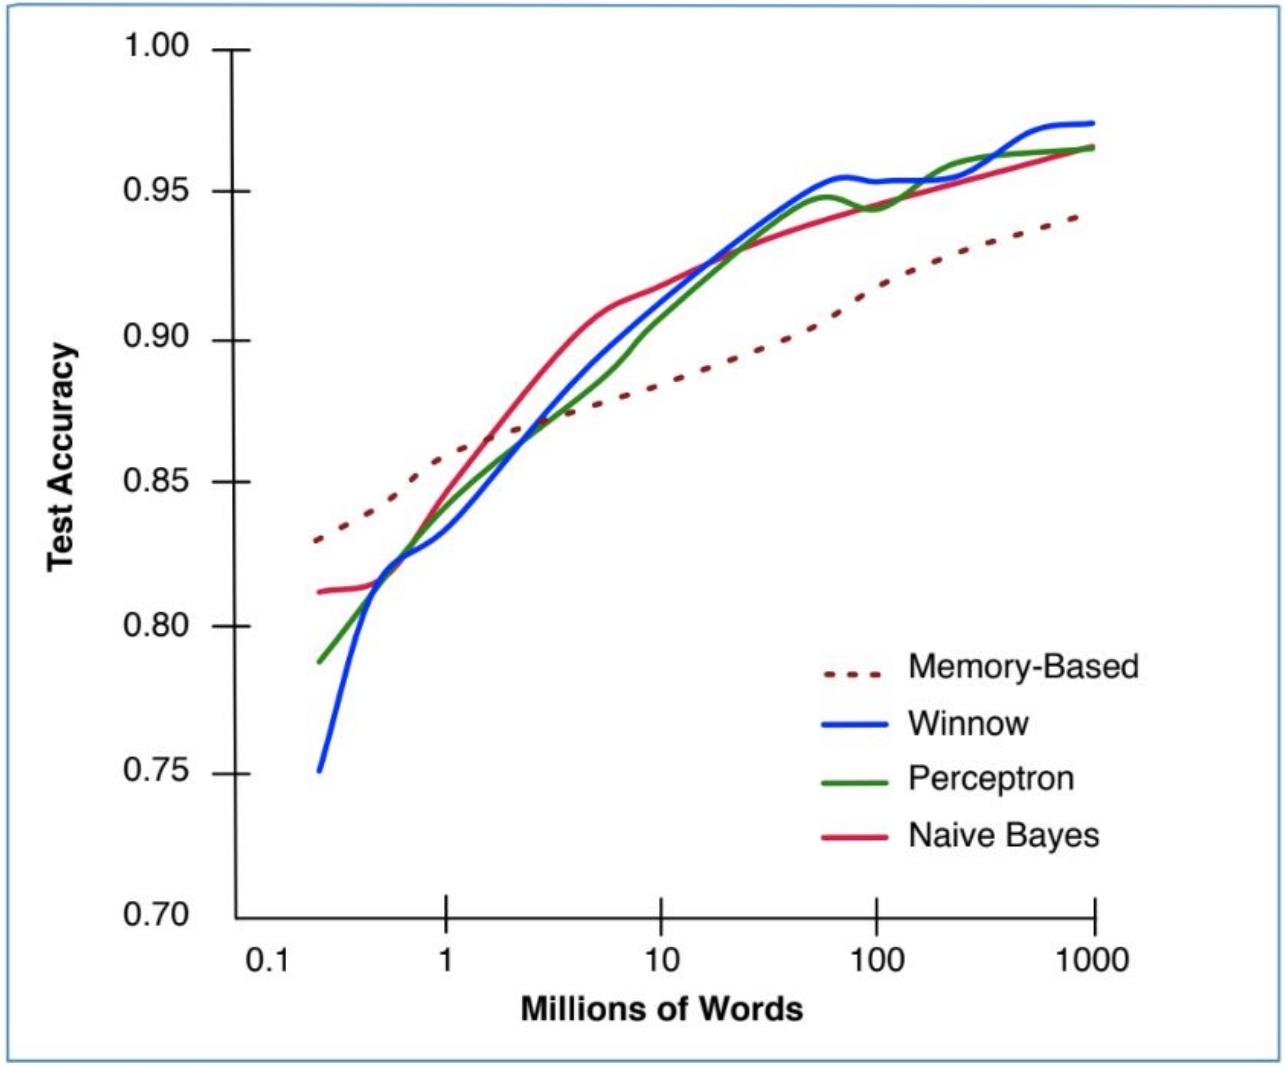

HOW IS THIS SOLVED OUTSIDE OF R?

HOW IS THIS SOLVED OUTSIDE OF R?

One of the many Jeff Dean facts:

Jeff Dean's PIN is the last 4 digits of pi.

The rest are funny-ish, but a good way to kill 20 mins on a many person webex.

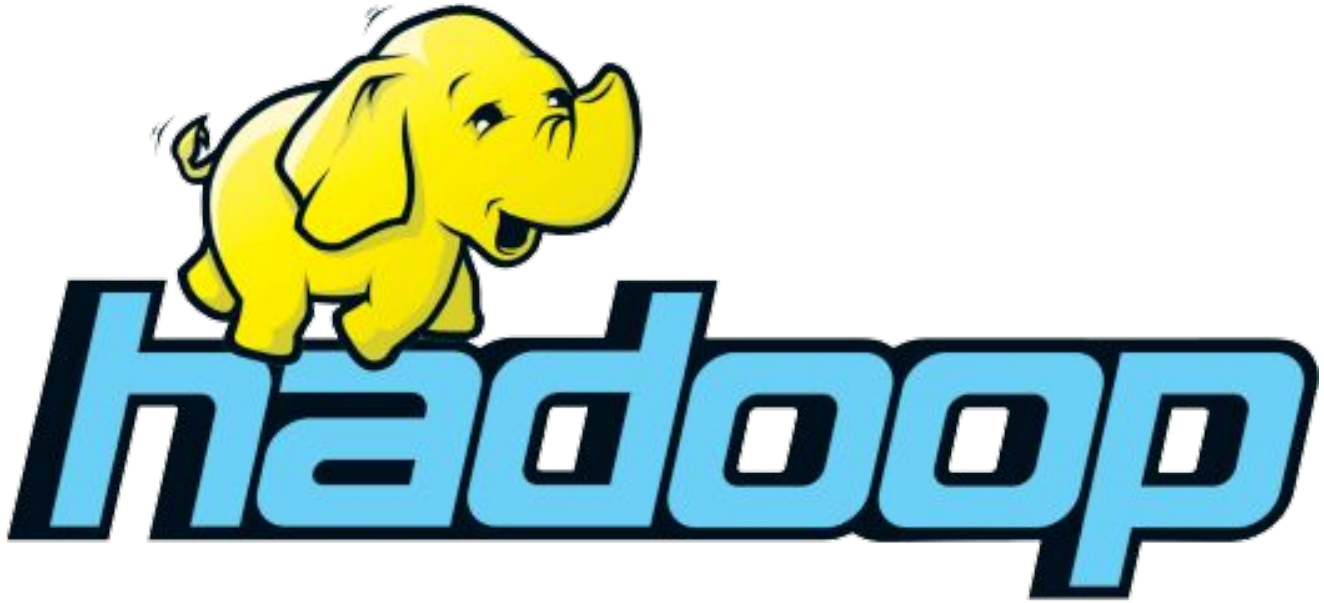

HOW IS THIS SOLVED OUTSIDE OF R?

Most of my colleagues who visit SA ask if they can eat some Impala. And Kudu. It's our other thing. African antelope are the spirit animals of database developers it seems.

A QUICK HADOOP OVERVIEW

HDFS

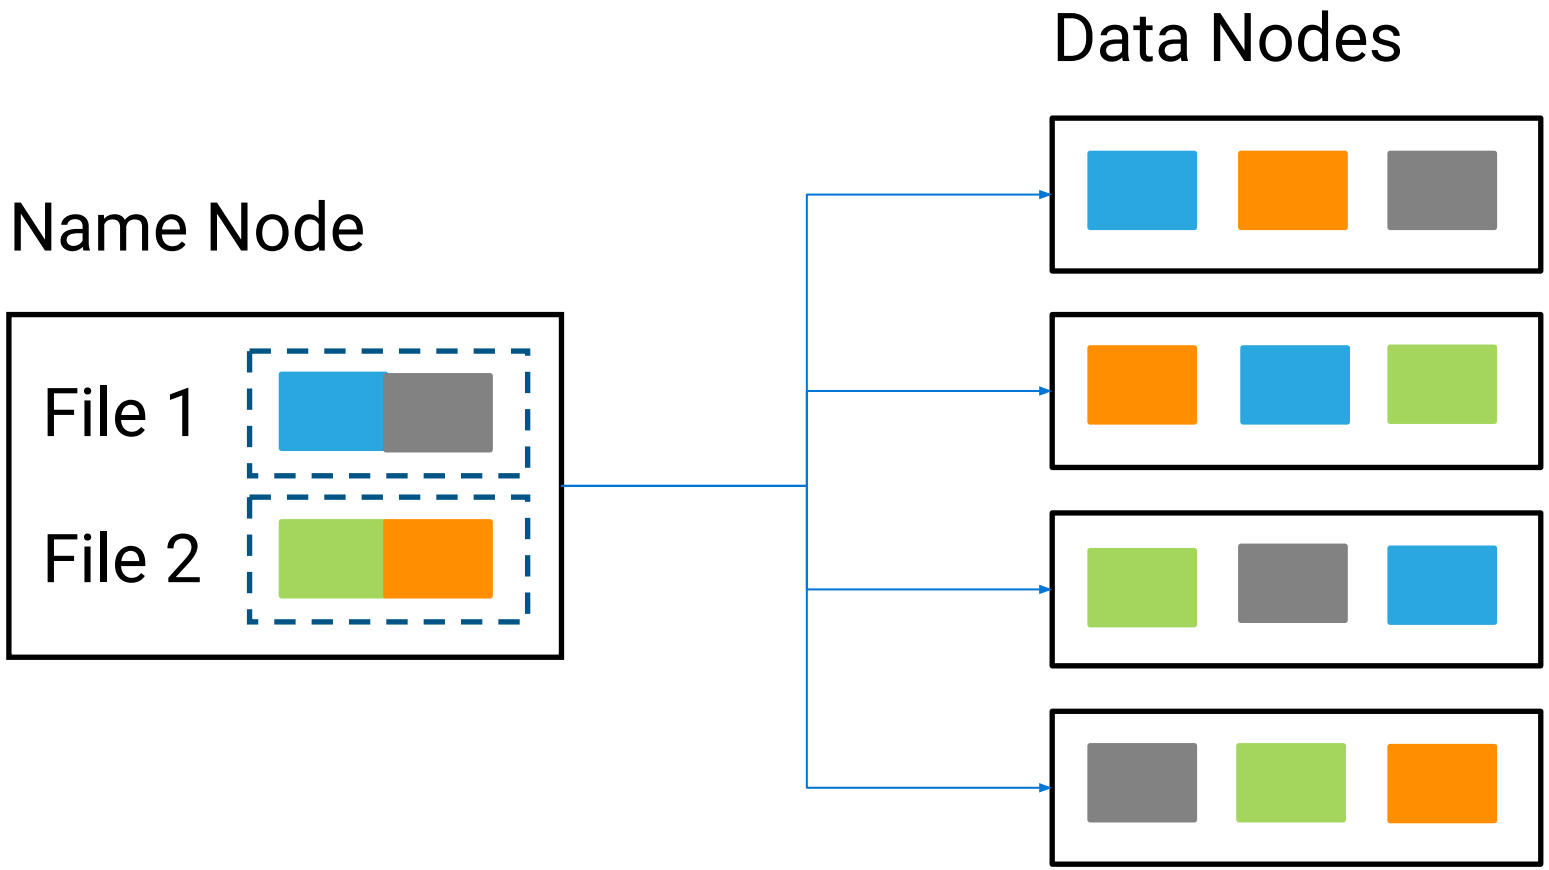

YARN

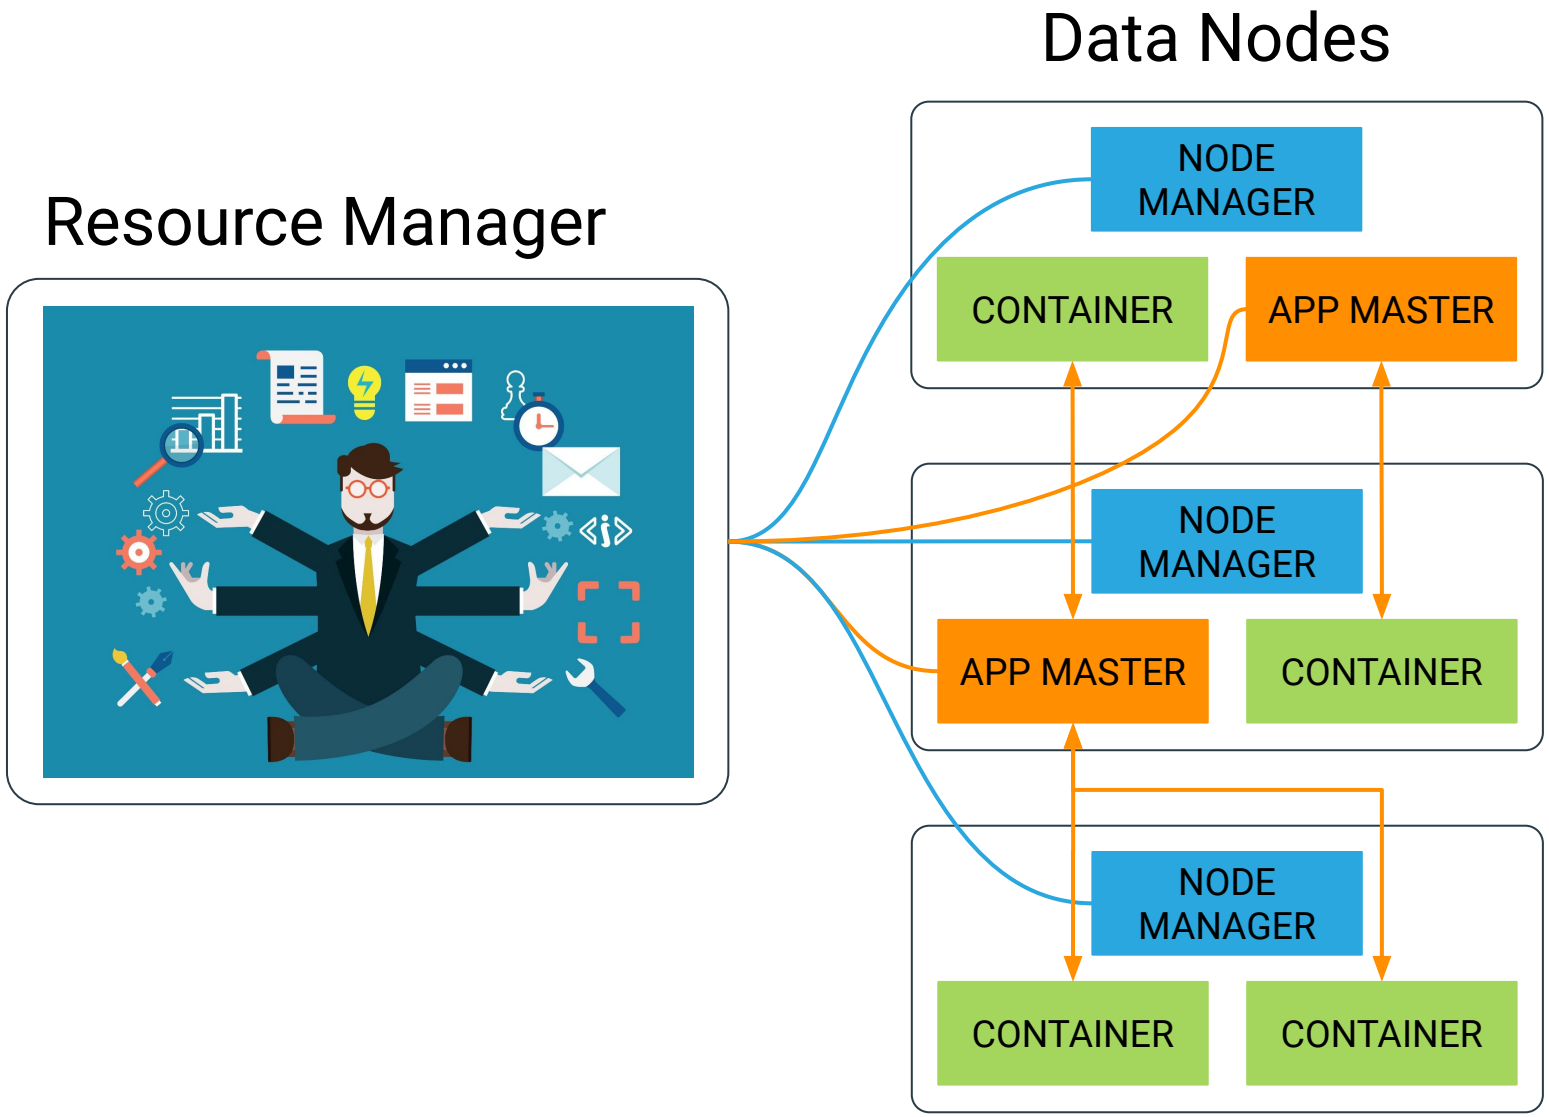

HADOOP STACK

Our template says I should put some related text here. How about - this layout was approved by Piet Mondrian.

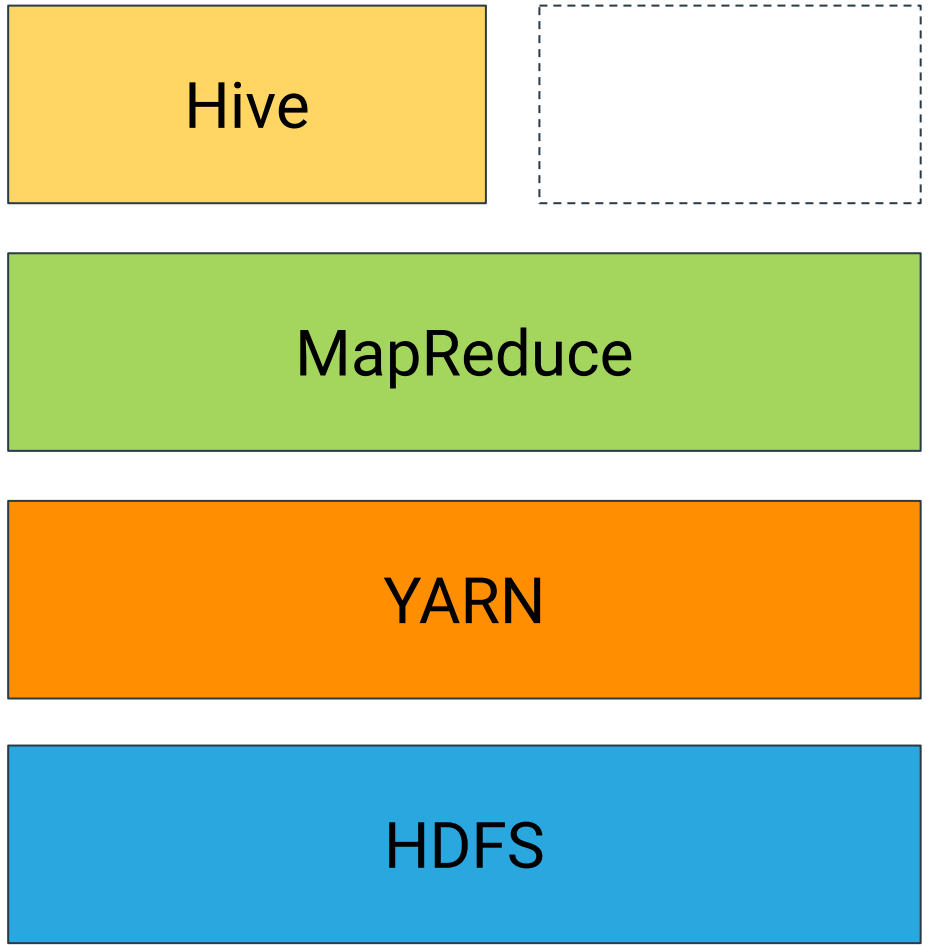

IMPALA OVERVIEW

HADOOP STACK WITH IMPALA

A Dutch painter and theoretician who is regarded as one of the greatest artists of the 20th century. Painty Piet I like to call him.

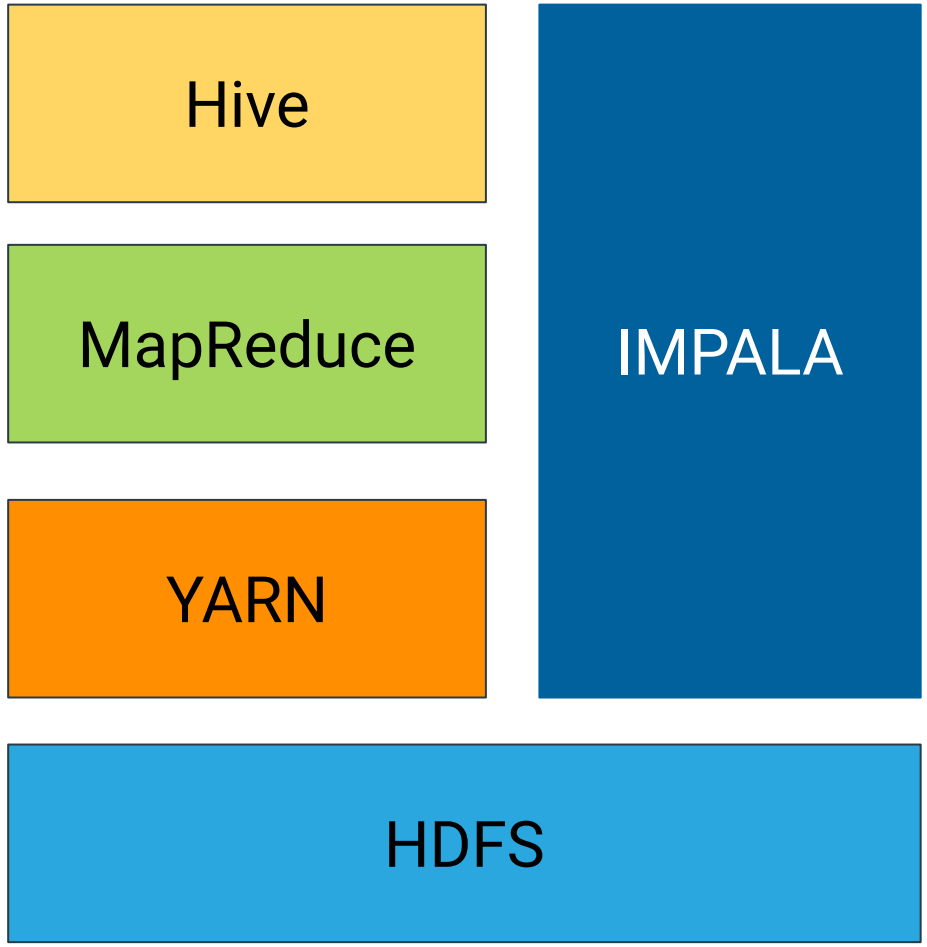

SPARK OVERVIEW

HADOOP STACK WITH SPARK

He proclaimed in 1914: Art is higher than reality and has no direct relation to reality…

According to Wikipedia anyway. I didn't hear him say it.

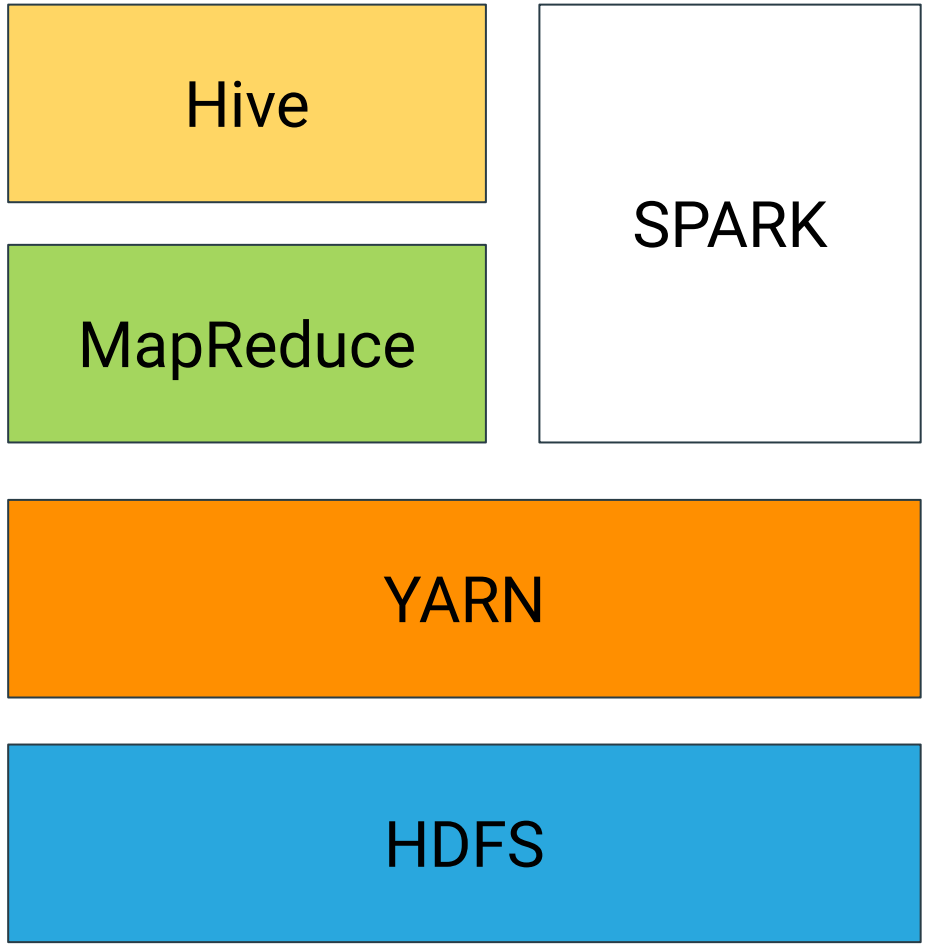

SPARK

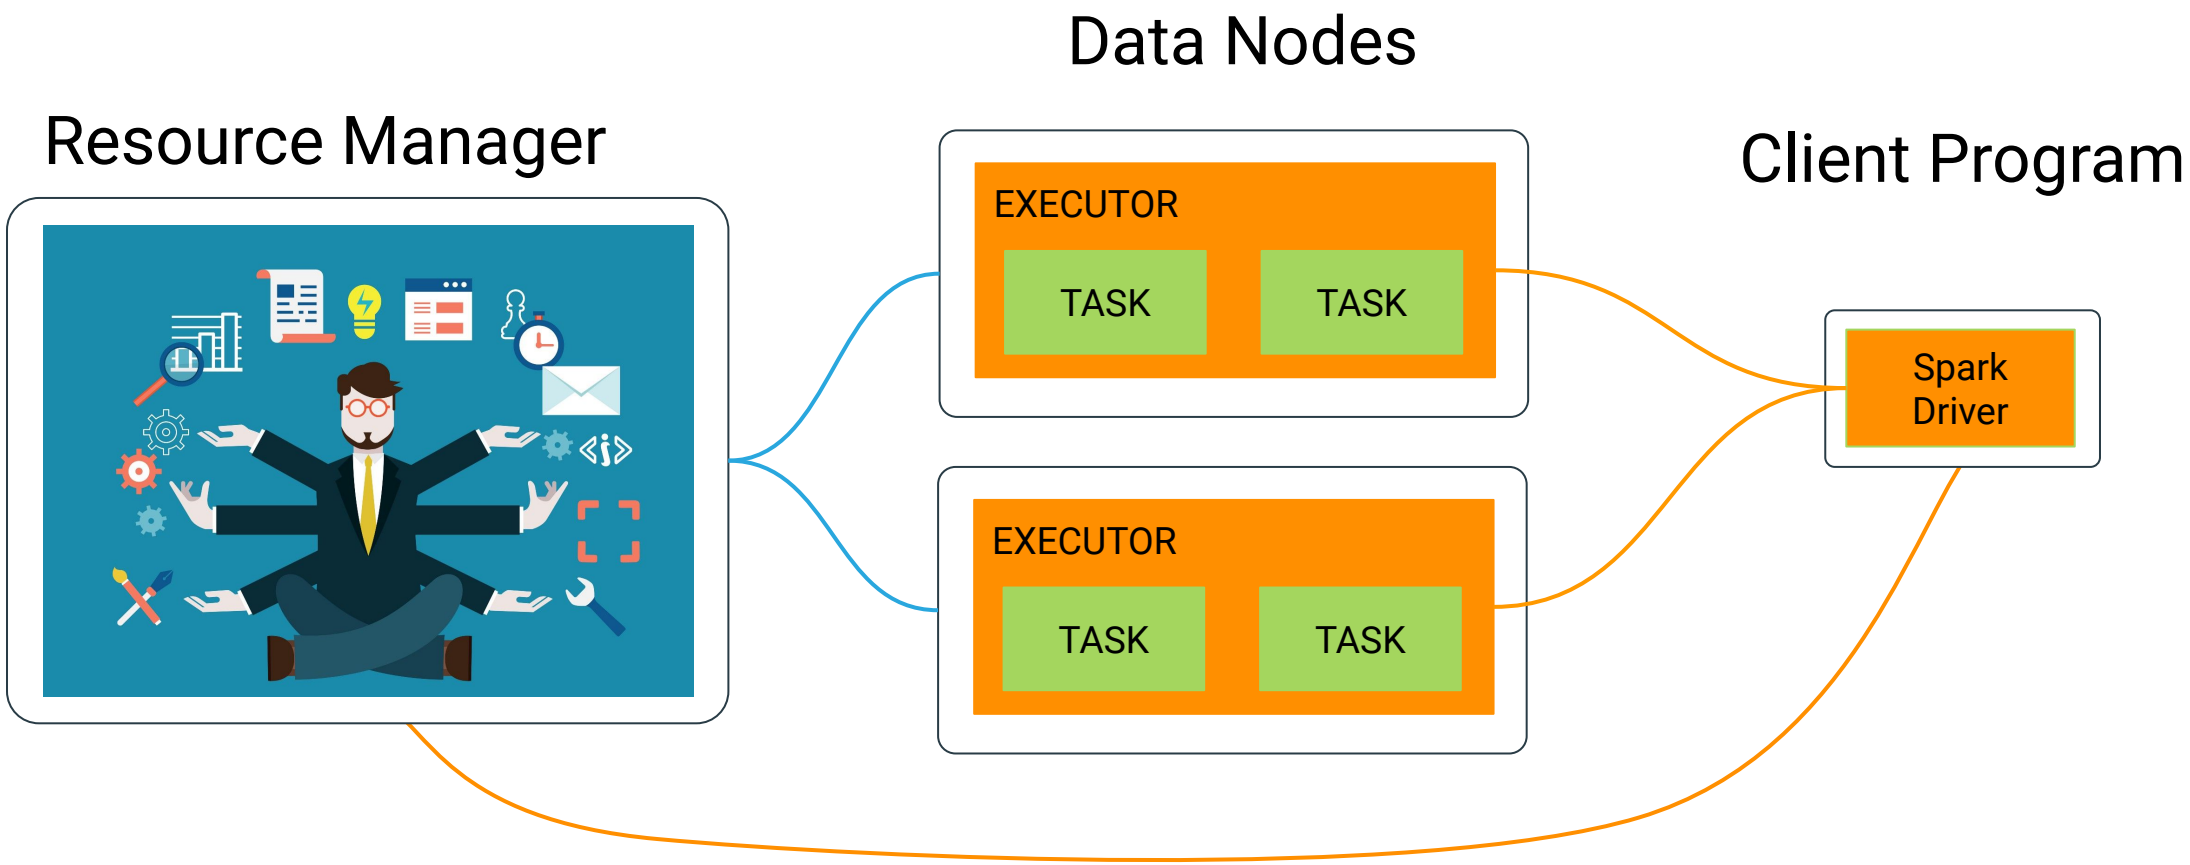

SPARK RDD

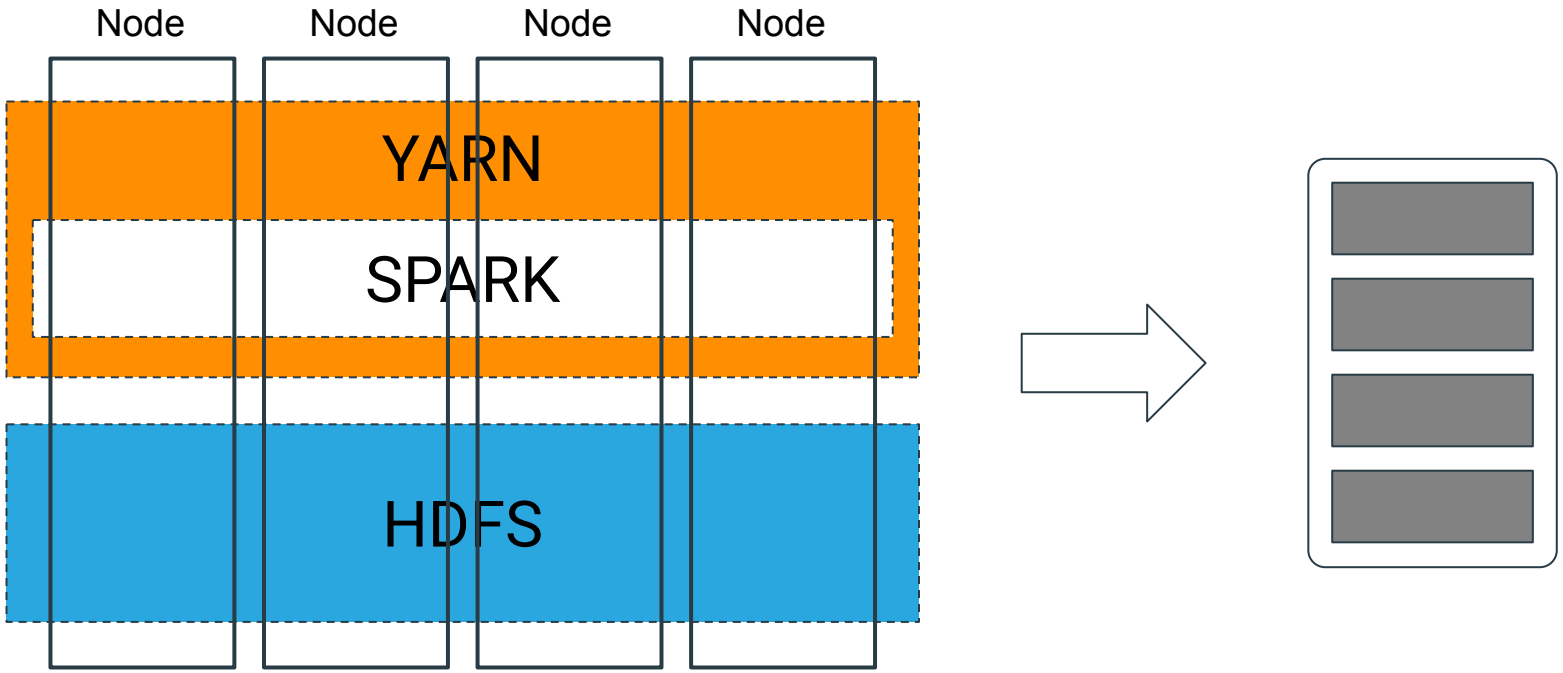

SPARK RDD

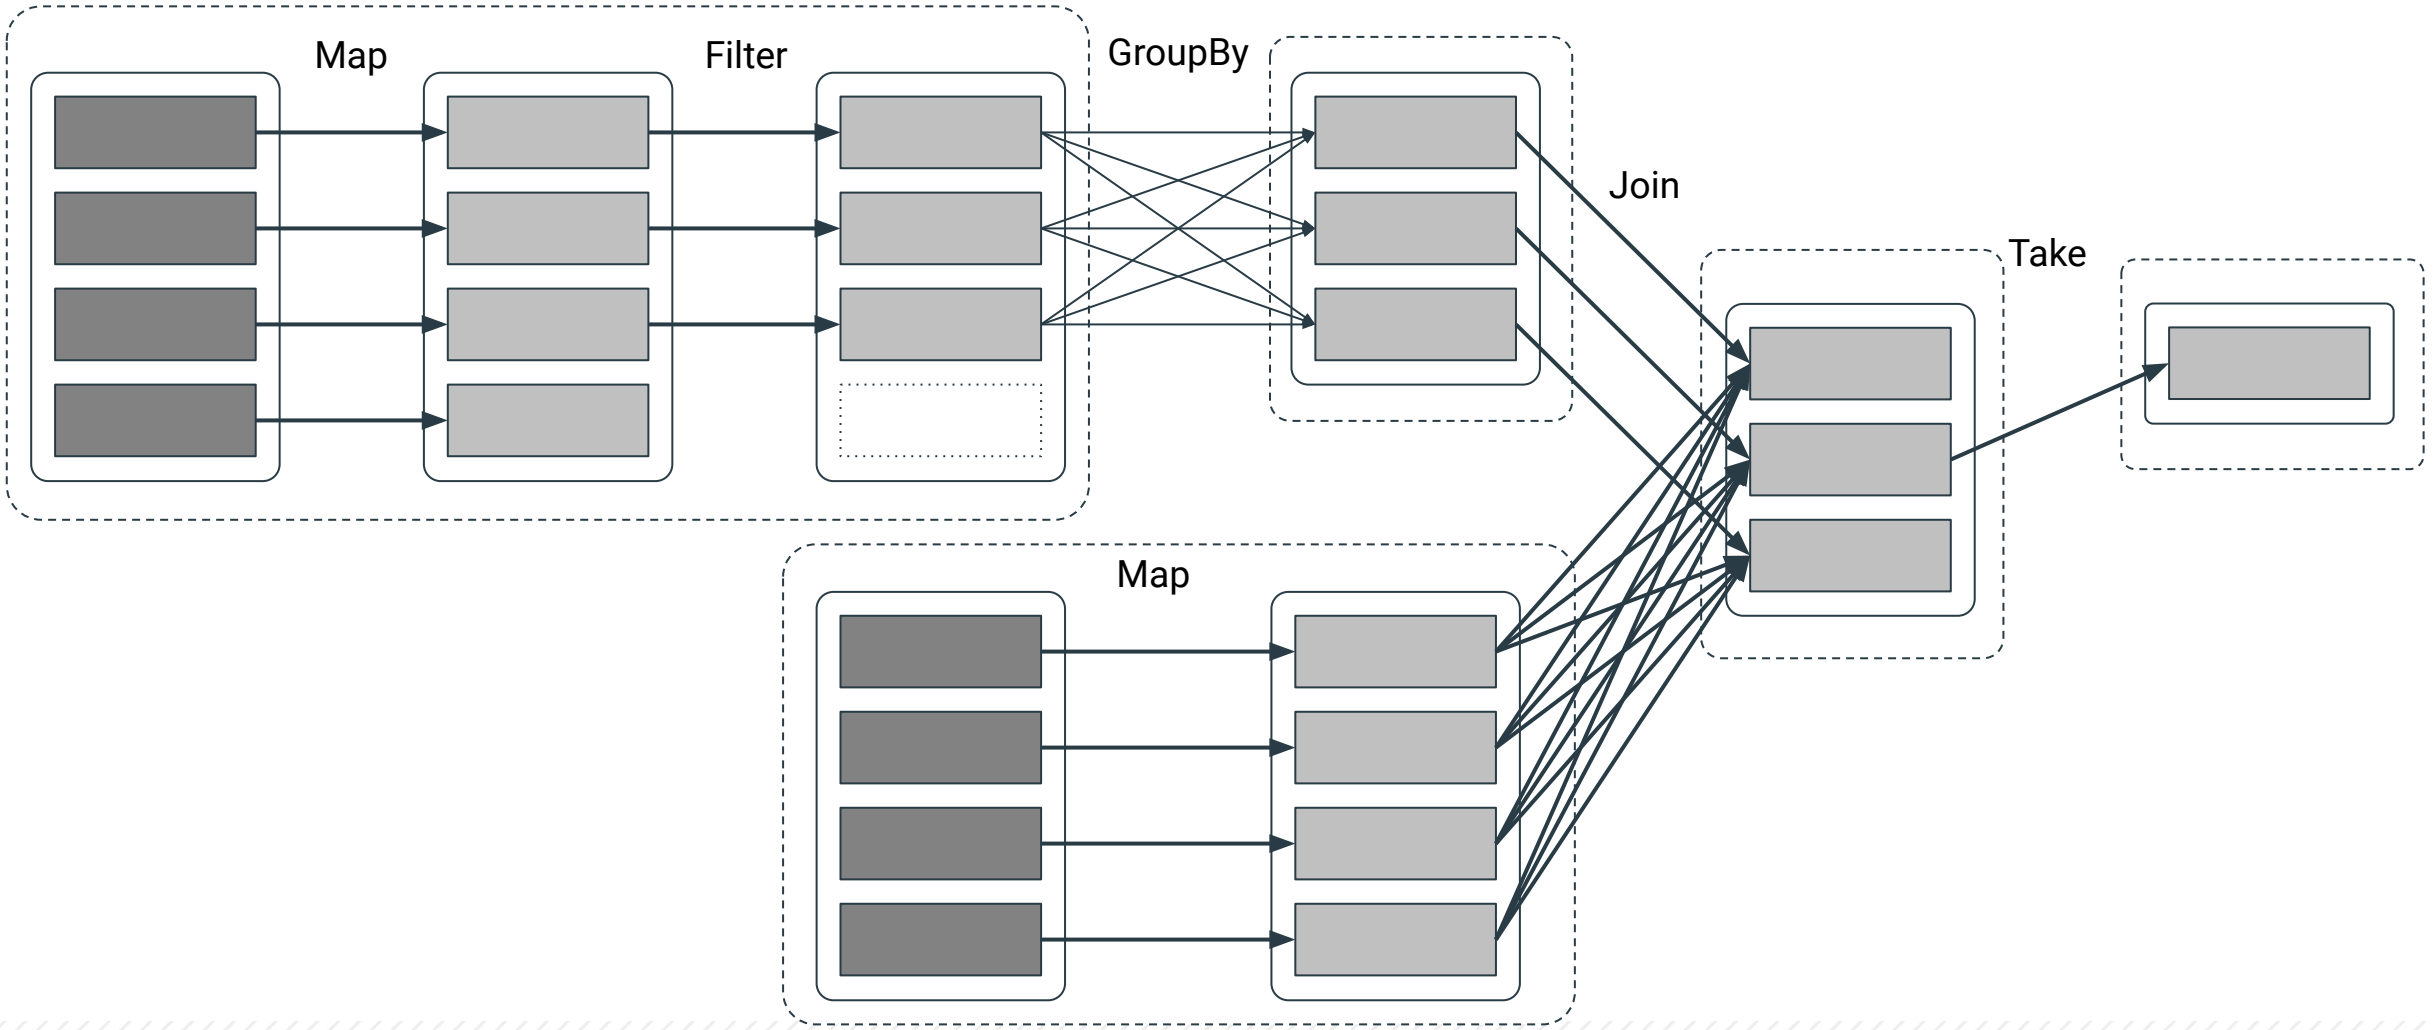

SPARK DATAFRAME

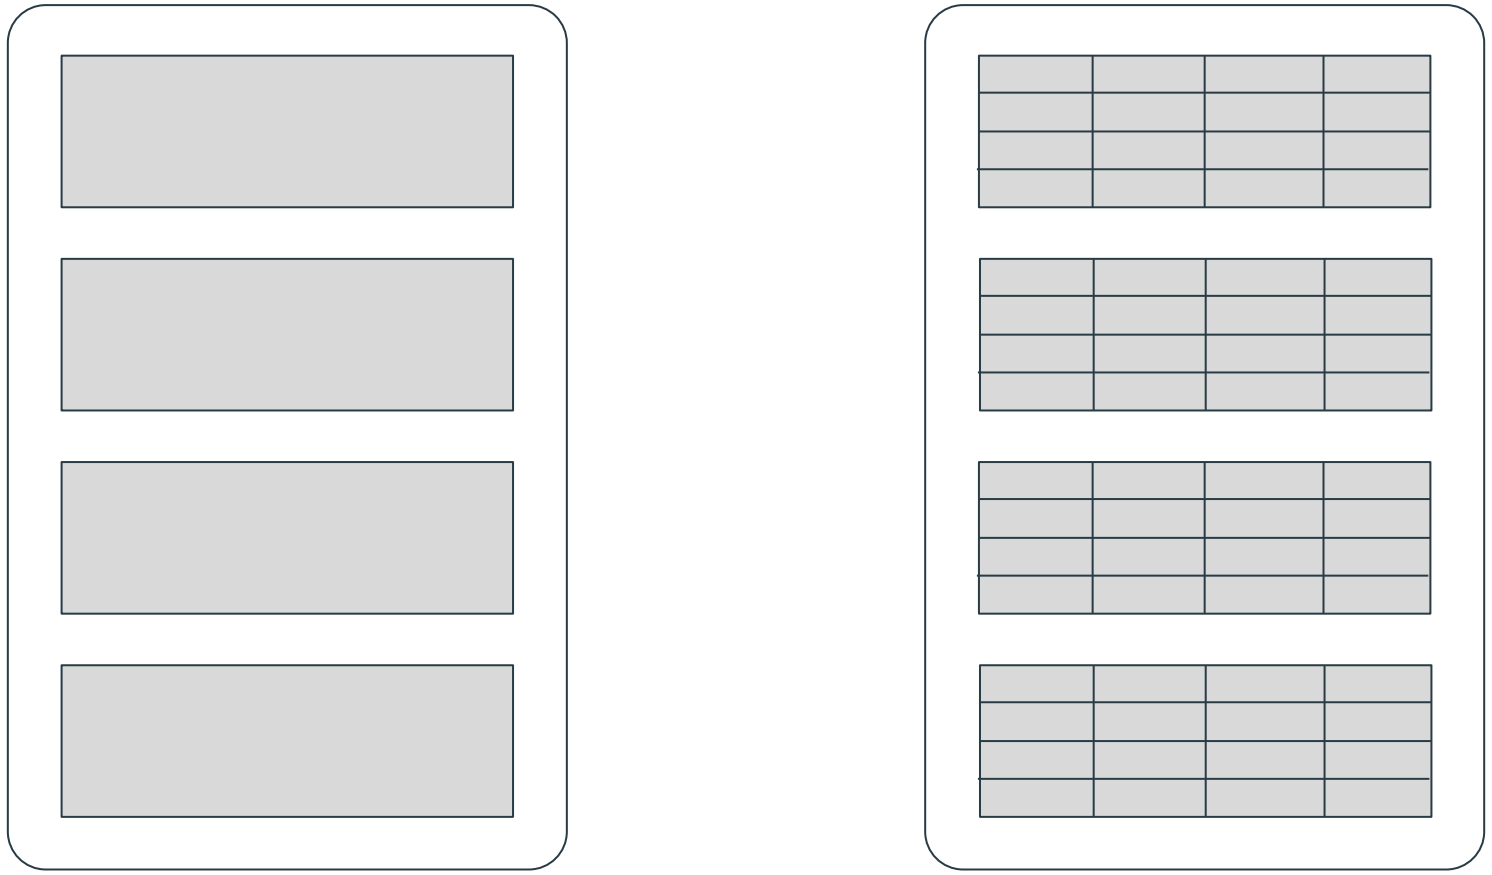

SPARK LIBRARIES

Apart from libraries, there are many books on Spark. Look for stuff by Holden Karau.

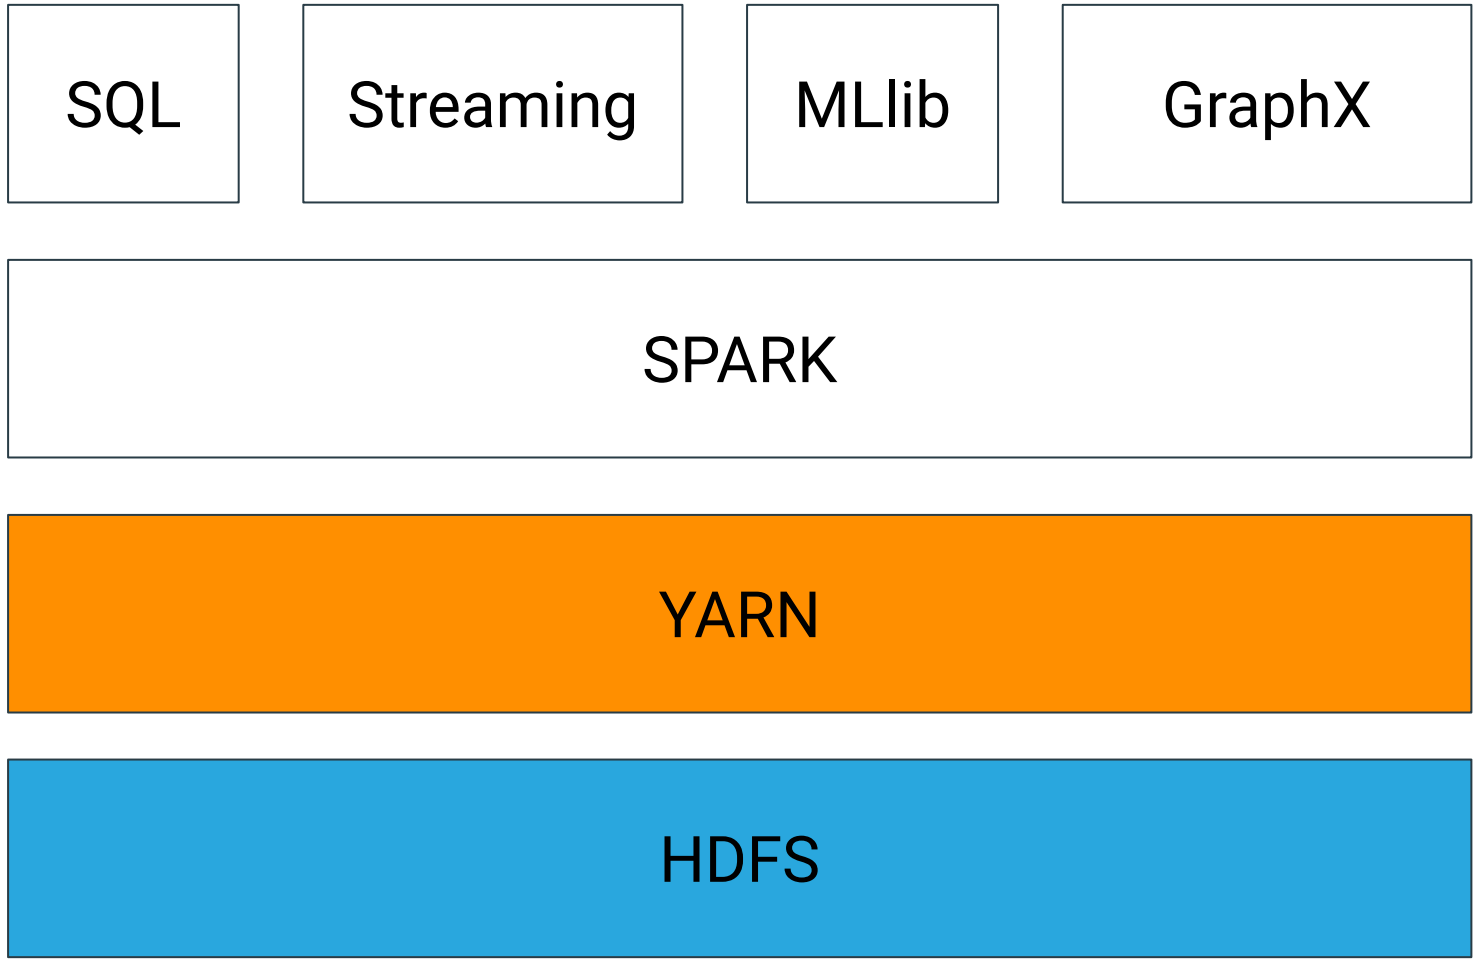

HOW DOES THIS WORK WITH R?

R vs Python

Scala wins, obviously!

HOW DOES THIS WORK WITH R?

implyr

```
delay <- flights_tbl %>%
  select(tailnum, distance, arr_delay) %>%
  group_by(tailnum) %>%
  summarise(count = n(), dist = mean(distance), delay = mean(arr_delay)) %>%
  filter(count > 20L, dist < 2000L, !is.na(delay)) %>%
  arrange(delay, dist, count) %>%
  collect()
```

SPARKLYR

The silver medal goes to SparkR. Sparklyr 1.0.0 has most of the things you need and support for Arrow.

SPARKLYR

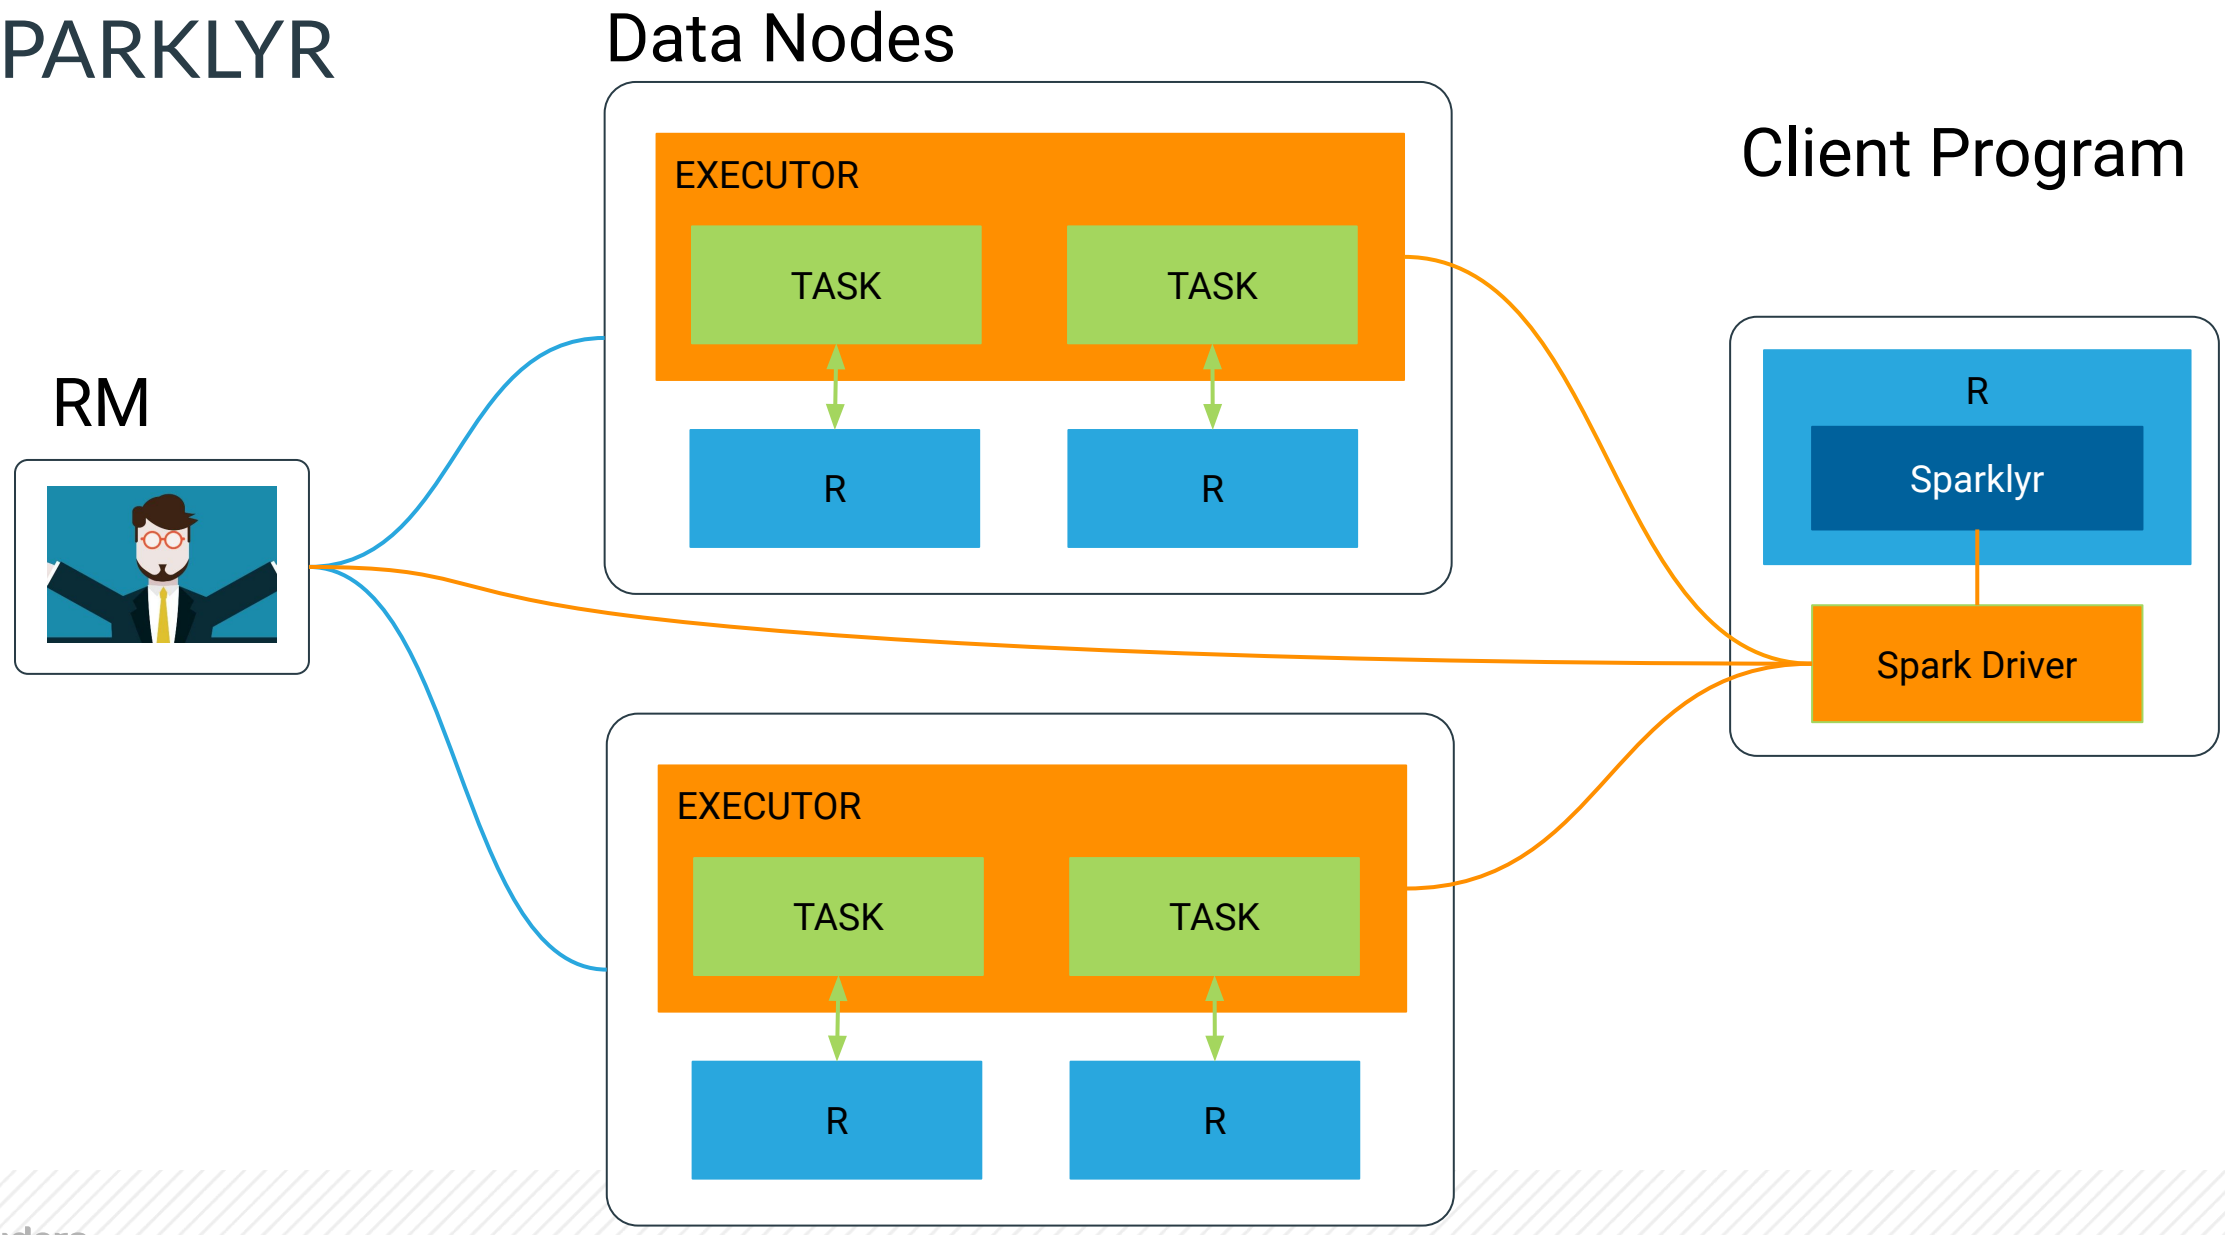

APACHE ARROW

High speed in memory data processing for everybody!

	session_id	timestamp	source_ip
Row 1	1331246660	3/8/2012 2:44PM	99.155.155.225
Row 2	1331246351	3/8/2012 2:38PM	65.87.165.114
Row 3	1331244570	3/8/2012 2:09PM	71.10.106.181
Row 4	1331261196	3/8/2012 6:46PM	76.102.156.138

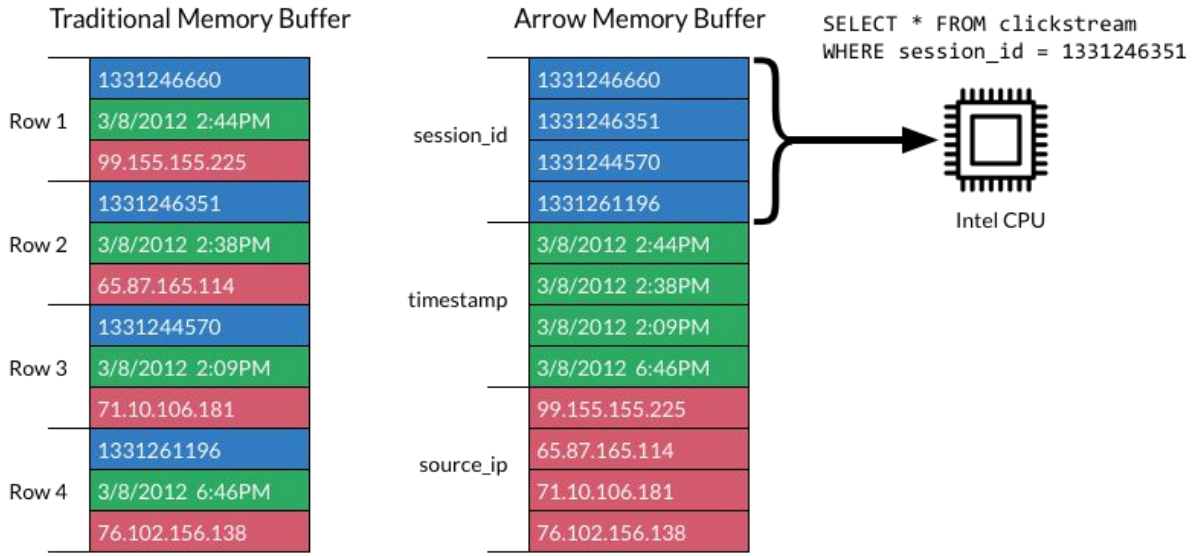

USEFUL THINGS TO KNOW

LAZY EVALUATION

This catches people out when working on the command line. Entered does not mean executed.

OVERLOADING THE DRIVER

I see this often when talking with data science teams. It's part of the standard operating procedure for "small data" data science work.

DATAFLOW

We’re moving into data engineering territory here but the pay is good and coffee is fair trade, locally sourced.

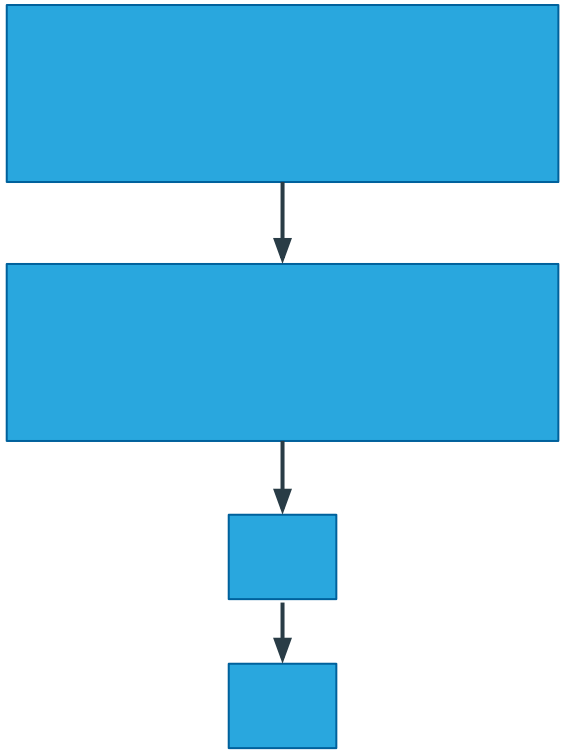

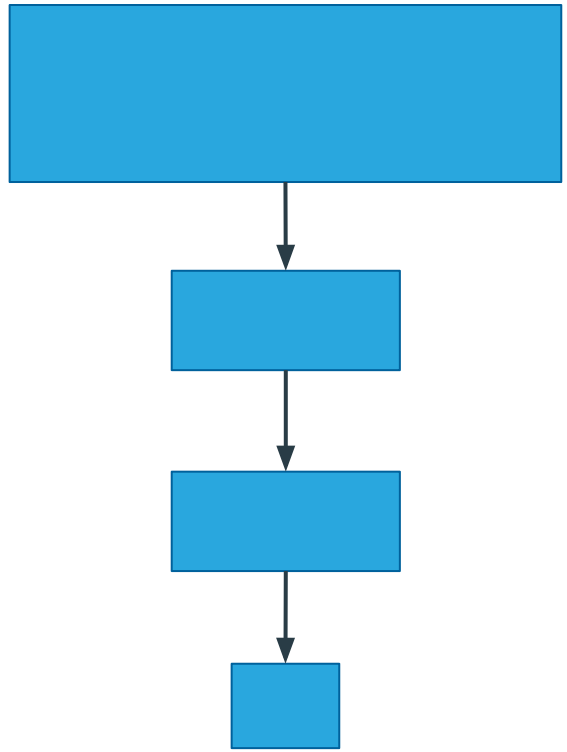

AVOID GROUPBYKEY

Use ReduceByKey. All the cool kids are doing it.

CHECK FOR DATA SKEW

I did not spend a lot of time drawing this. But Edward Tufte would be proud. Very favourable data to ink ratio.

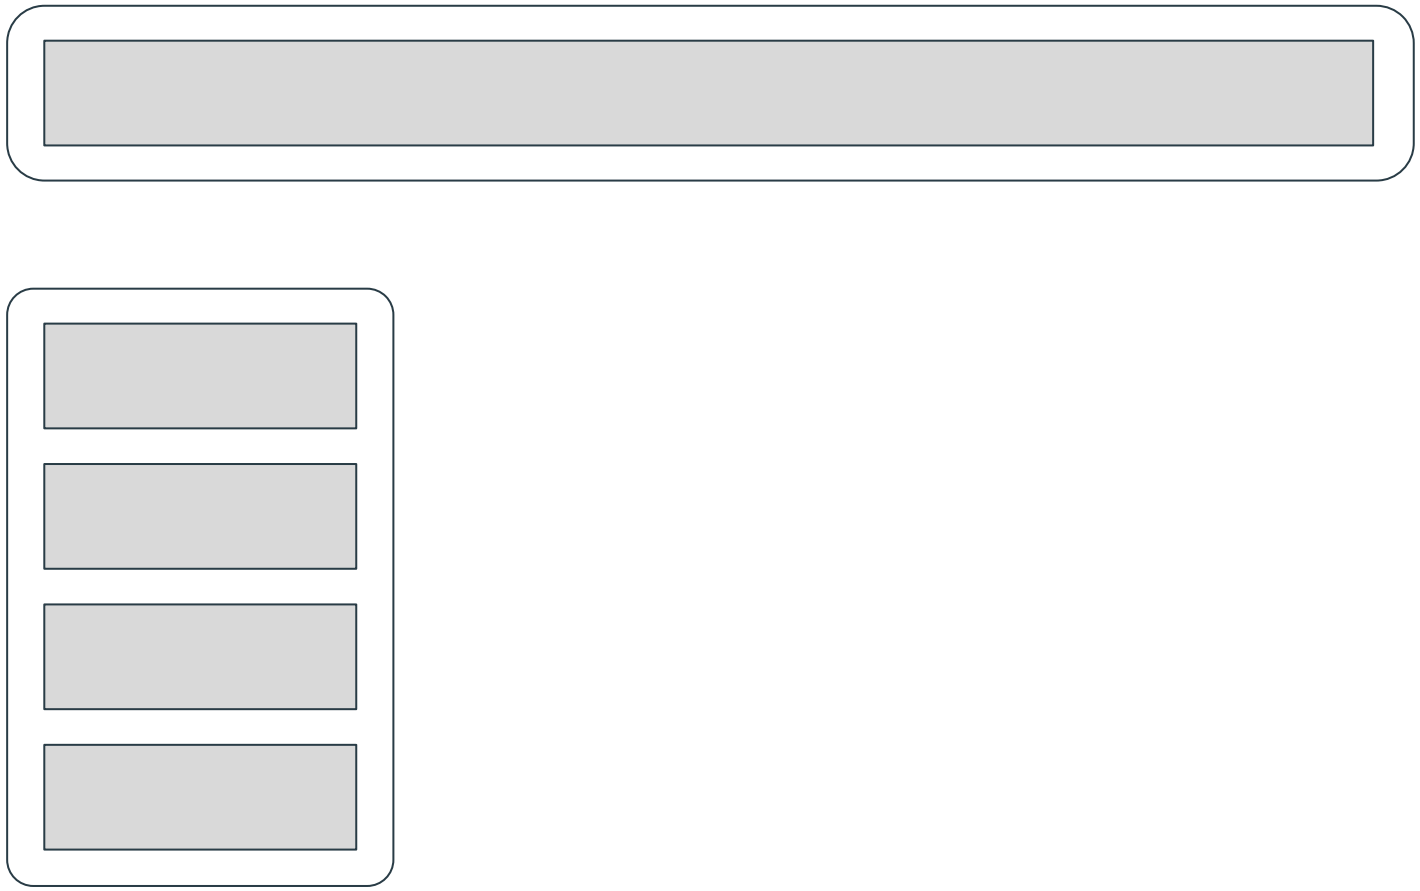

CHECK FOR DATA SKEW

You can avoid this by adding salt to the keys. Definitely no easy set up for a joke with that.

USE THE SPARK UI

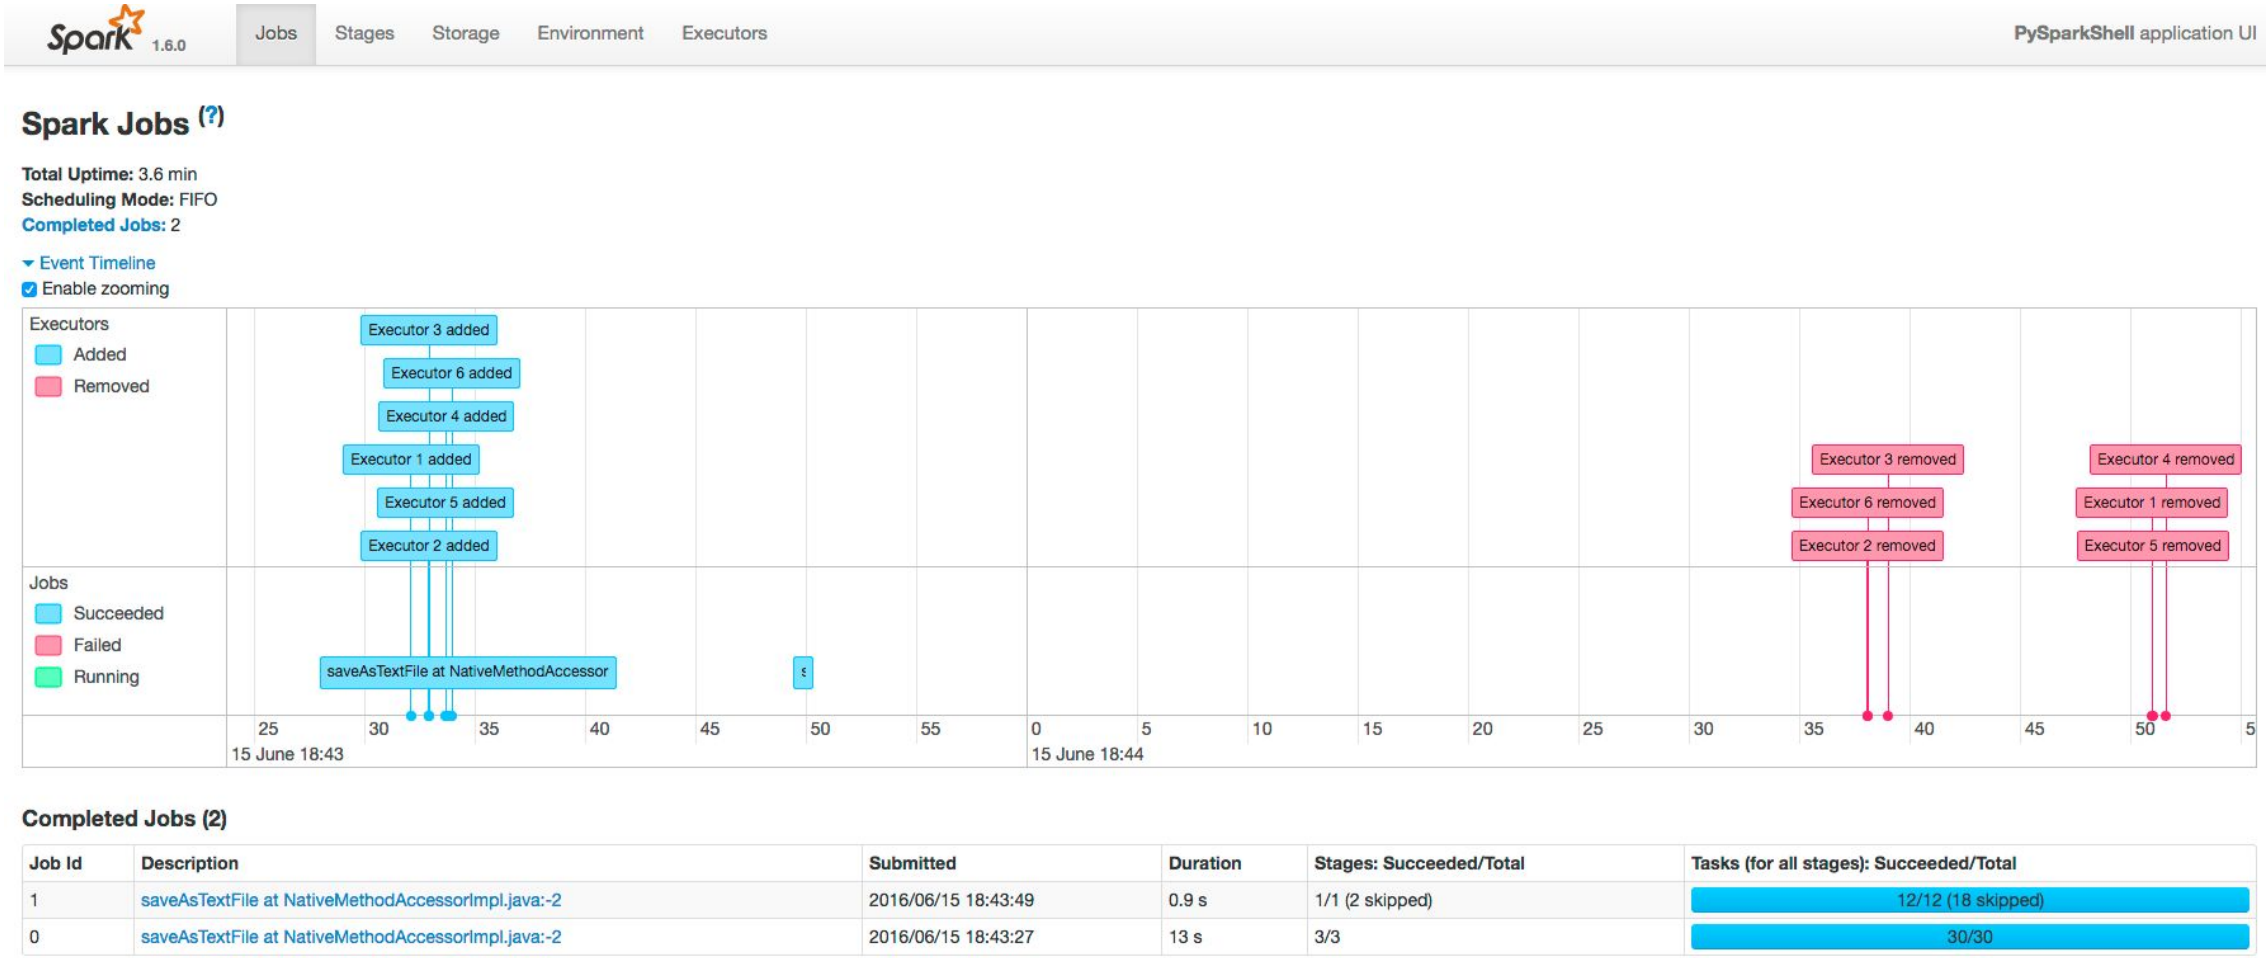

Completed Jobs (2)

Job Id	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	saveAsTextFile at NativeMethodAccessorImpl.java:-2	2016/06/15 18:43:49	0.9 s	1/1 (2 skipped)	12/12 (18 skipped)
0	saveAsTextFile at NativeMethodAccessorImpl.java:-2	2016/06/15 18:43:27	13 s	3/3	30/30

DAG

Directed Acyclic Graph -
Use the full title if you want to sound smarter while talking data scientists.

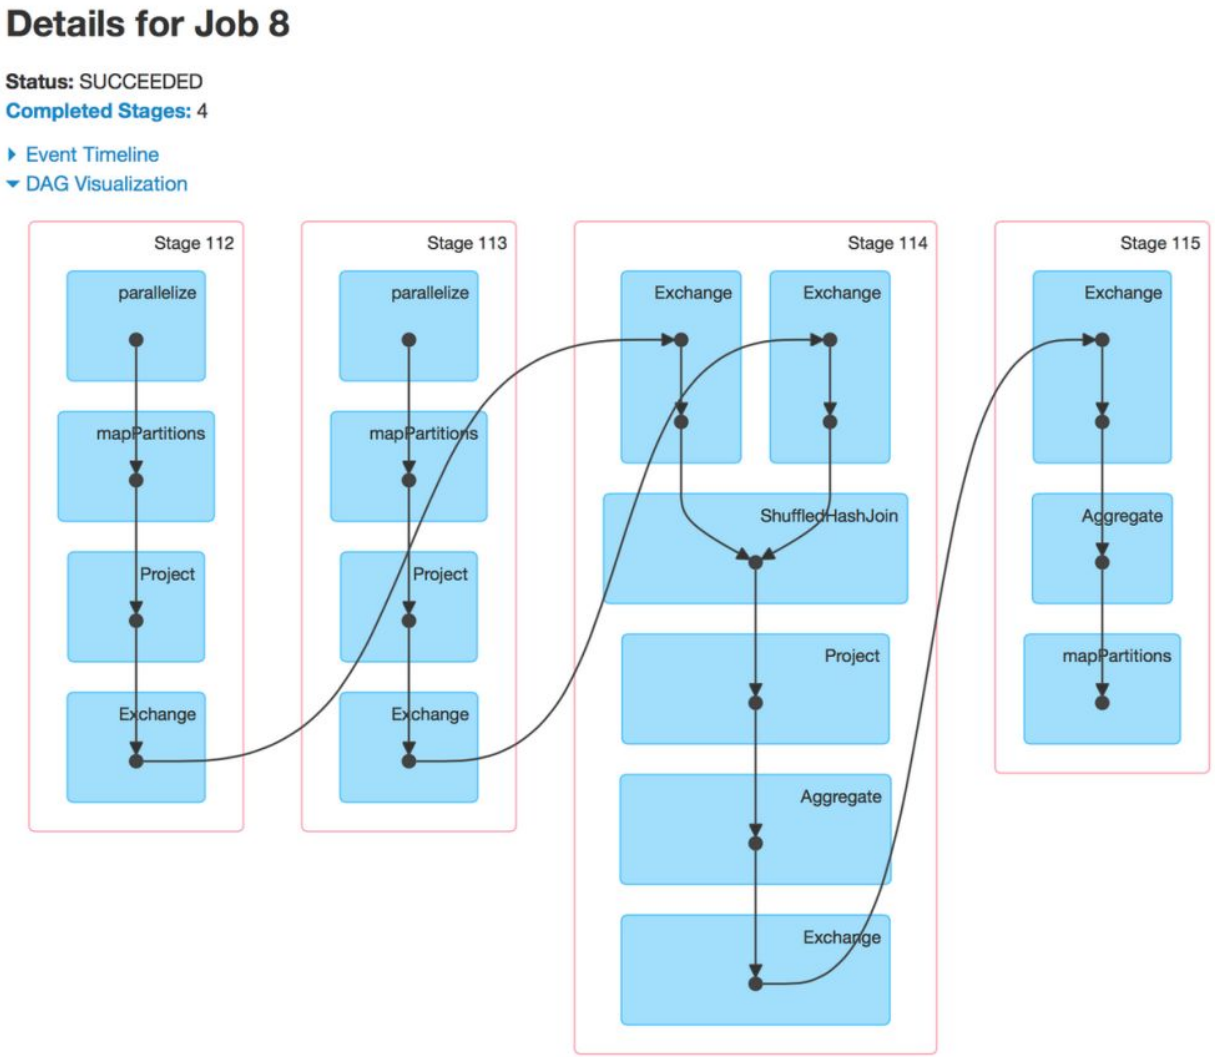

GETTING STARTED

GETTING STARTED

Implyr

https://blog.cloudera.com/blog/2017/07/implyr-r-interface-for-apache-impala/

https://github.com/ianmcook/implyr

Sparklyr

https://blog.cloudera.com/blog/2016/09/introducing-sparklyr-an-r-interface-for-apache-spark/

https://spark.rstudio.com/

Spark

https://www.cloudera.com/documentation/enterprise/5-16-x/PDF/cloudera-spark.pdf

https://www.cloudera.com/products/open-source/apache-hadoop/apache-spark.html

THANK YOU

IMAGES USED

Most diagrams I either drew or took from Cloudera's documentation directly. The rest come from here:

https://www.apple.com/shop/buy-mac/macbook-pro/15-inch-space-gray-2.6ghz-6-core-512gb

http://web.cs.ucla.edu/classes/winter13/cs111/scribe/10c/

https://commons.wikimedia.org/wiki/File:Special_Edition_NYC_2015_-_DC_vs_Street_Fighter_(18357923460).jpg

https://commons.wikimedia.org/wiki/File:Overloaded_truck.jpg

https://www.publicdomainpictures.net/en/view-image.php?image=220900&picture=two-lazy-bears

https://www.goodfreephotos.com/other-photos/lots-of-keys.jpg.php